

Problem Statement:

An education department in the US needs to analyse the factors that influence the admission of a student into a college.

Analyse the historical data and determine the key drivers.

Analysis information:

Predictive

- Run logistic model to determine the factors that influence the admission process of a student (Drop insignificant variables)
- Transform variables to factors wherever required
- Calculate accuracy of the model
- Try other modeling techniques like decision tree and SVM and select a champion model
- Determine the accuracy rates for each model
- Select the most accurate model
- Identify other Machine learning or statistical techniques that can be used

Descriptive

- Categorize the grade point average into High, Medium, and Low (with admission probability percentages) and plot it on a point chart.
- Cross grid for admission variables with GRE Categorization is shown below:

GRE	Categori zed
0-440	Low
440-580	Medium
580 +	High

College Admission Prediction

Exploring the dataset:

```
Console Terminal x Jobs x
~/
> head(col_adm)
  admit gre  gpa ses Gender_Male Race rank
1     0 380 3.61  1          0     3    3
2     1 660 3.67  2          0     2    3
3     1 800 4.00  2          0     2    1
4     1 640 3.19  1          1     2    4
5     0 520 2.93  3          1     2    4
6     1 760 3.00  2          1     1    2
> # No of records
> nrow(col_adm)
[1] 400
> # Count of missing values
> sum(is.na(col_adm))
[1] 0
> # Name of the columns
> names(col_adm)
[1] "admit"      "gre"        "gpa"        "ses"
[5] "Gender_Male" "Race"       "rank"
> summary(col_adm)
      admit      gre      gpa
Min.   :0.0000   Min.   :220.0   Min.   :2.260
1st Qu.:0.0000   1st Qu.:520.0   1st Qu.:3.130
Median :0.0000   Median :580.0   Median :3.395
Mean   :0.3175   Mean   :587.7   Mean   :3.390
3rd Qu.:1.0000   3rd Qu.:660.0   3rd Qu.:3.670
Max.   :1.0000   Max.   :800.0   Max.   :4.000
      ses      Gender_Male      Race
Min.   :1.000   Min.   :0.000   Min.   :1.000
1st Qu.:1.000   1st Qu.:0.000   1st Qu.:1.000
Median :2.000   Median :0.000   Median :2.000
Mean   :1.992   Mean   :0.475   Mean   :1.962
3rd Qu.:3.000   3rd Qu.:1.000   3rd Qu.:3.000
Max.   :3.000   Max.   :1.000   Max.   :3.000
      rank
Min.   :1.000
1st Qu.:2.000
Median :2.000
Mean   :2.485
3rd Qu.:3.000
Max.   :4.000
```

```

> # Data type of columns
> sapply(col_adm,class)
      admit      gre      gpa      ses Gender_Male
"integer" "integer" "numeric" "integer" "integer"
      Race      rank
"integer" "integer"
> #Converting columns to factor type
> col_adm$admit=as.factor(col_adm$admit)
> col_adm$ses=as.factor(col_adm$ses)
> col_adm$Gender_Male=as.factor(col_adm$Gender_Male)
> col_adm$Race=as.factor(col_adm$Race)
> col_adm$rank=as.factor(col_adm$rank)
> # Checking the datatype after converting :
> sapply(col_adm,class)
      admit      gre      gpa      ses Gender_Male
"factor"  "integer" "numeric" "factor" "factor"
      Race      rank
"factor"  "factor"

```

Outlier Detection and Removal:

```

> #outlier detection
> #for gre
>
> iqr1=IQR(col_adm$gre)
> iqr1
[1] 140
> quantile(col_adm$gre,na.rm = T)
 0%  25%  50%  75% 100%
220 520 580 660 800
> max1=660+1.5*iqr1
> max1
[1] 870
> min1=520-1.5*iqr1
> min1
[1] 310
>
> # All the points above the upperInner fence
> print(which(col_adm$gre>max1)) # no outlier
integer(0)
> print(which(col_adm$gre<min1)) # 4 outliers
[1] 72 180 305 316
> # for gpa variable
> iqr2=IQR(col_adm$gpa)
> iqr2
[1] 0.54
> quantile(col_adm$gpa,na.rm = T)
 0%  25%  50%  75% 100%
2.260 3.130 3.395 3.670 4.000
> max2=3.670+1.5*iqr2
> min2=3.130-1.5*iqr2
>
> # All the points above the upperInner fence
> print(which(col_adm$gpa>max2)) ## No outlier
integer(0)
>
> #all the points below the lowerInner fence
> print(which(col_adm$gpa<min2)) # 1 outlier
[1] 290
>

```

```
> # Removal of outlier
>
> col_adm=col_adm[-c(72,180,305,316,290),]
> nrow(col_adm) # Outlier removed
[1] 395
> |
```

Data Splitting:

```
> # Splitting the data into train and test
>
> set.seed(0)
> library('caTools')
> col_adm[,c(2,3)]=scale(col_adm[,c(2,3)])
> split=sample.split(col_adm$admit,SplitRatio = 0.75)
> train=subset(col_adm,split==T)
> test=subset(col_adm,split==F)
> |
```

Logistic Regression:

```
> # Logistic Regression
>
> logit1=glm(admit~.,train,family = 'binomial')
> summary(logit1)
```

Call:

```
glm(formula = admit ~ ., family = "binomial", data = train)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-1.8458	-0.8294	-0.5794	0.9459	2.1850

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.03714	0.45482	2.280	0.02259 *
gre	0.27452	0.15262	1.799	0.07206 .
gpa	0.51793	0.16125	3.212	0.00132 **
ses2	-0.38459	0.33468	-1.149	0.25050
ses3	-0.40768	0.34356	-1.187	0.23538
Gender_Male1	-0.09887	0.27541	-0.359	0.71959
Race2	-0.34389	0.33981	-1.012	0.31153
Race3	-0.43592	0.33303	-1.309	0.19054
rank2	-1.29613	0.40854	-3.173	0.00151 **
rank3	-1.70399	0.43413	-3.925	8.67e-05 ***
rank4	-2.05159	0.51218	-4.006	6.19e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 370.01 on 295 degrees of freedom
 Residual deviance: 319.50 on 285 degrees of freedom
 AIC: 341.5

Number of Fisher Scoring iterations: 4

```
> |
```

Second Logistic model by removing the insignificant variable:

```
~/
> #So here gre,ses,gender_male and race variable are not significant.
> # So building new model with gpa and rank variable.
>
> logit2=glm(admit~gpa+rank,train,family = 'binomial')
> summary(logit2)
```

Call:

```
glm(formula = admit ~ gpa + rank, family = "binomial", data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.8203	-0.8527	-0.5997	1.0019	2.2480

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.4765	0.3386	1.407	0.15931
gpa	0.6037	0.1479	4.081	4.49e-05 ***
rank2	-1.2261	0.3978	-3.082	0.00205 **
rank3	-1.7363	0.4248	-4.087	4.37e-05 ***
rank4	-2.0552	0.5038	-4.079	4.52e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 370.01 on 295 degrees of freedom
Residual deviance: 326.90 on 291 degrees of freedom
AIC: 336.9

Number of Fisher Scoring iterations: 4

Accuracy of the logistic model:

```
> #Accuracy of Logistic Model
>
> predicted_val1=predict(logit1,test,type='response')
>
> test$pred_admit1=ifelse(predicted_val1>0.5,1,0)
>
> #creating the confusion matrix
>
> conf_mat1=table(predicted=test$pred_admit1,actual=test$admit)
> conf_mat1
      actual
predicted 0  1
      0 55 26
      1 12  6
>
> # Getting the accuracy
>
> accuracy1=sum(diag(conf_mat1))/sum(conf_mat1)
>
> ## Accuracy is 0.6161616
> |
```

He

SVM MODEL:

```
> #SVM Model
>
> library('e1071')
>
> svm_clf=svm(admit~.,train,type='C-classification',kernel='linear')
> summary(svm_clf)

Call:
svm(formula = admit ~ ., data = train, type = "C-classification", kernel = "linear")

Parameters:
  SVM-Type:  C-classification
 SVM-Kernel: linear
      cost:  1

Number of Support Vectors:  202

( 94 108 )

Number of Classes:  2

Levels:
0 1
```

Accuracy of SVM model:

```
> #accuracy of SVM model
>
> predicted_val2=predict(svm_clf,test[-1])
>
> # Creating the confusion matrix
>
> conf_mat2=table(predicted=predicted_val2,actual=test$admit)
> conf_mat1
      actual
predicted 0  1
      0 55 26
      1 12  6

>
> accuracy2=sum(diag(conf_mat2))/sum(conf_mat2)
> accuracy2
[1] 0.6161616
>
> # Accuracy is 0.61
> |
```

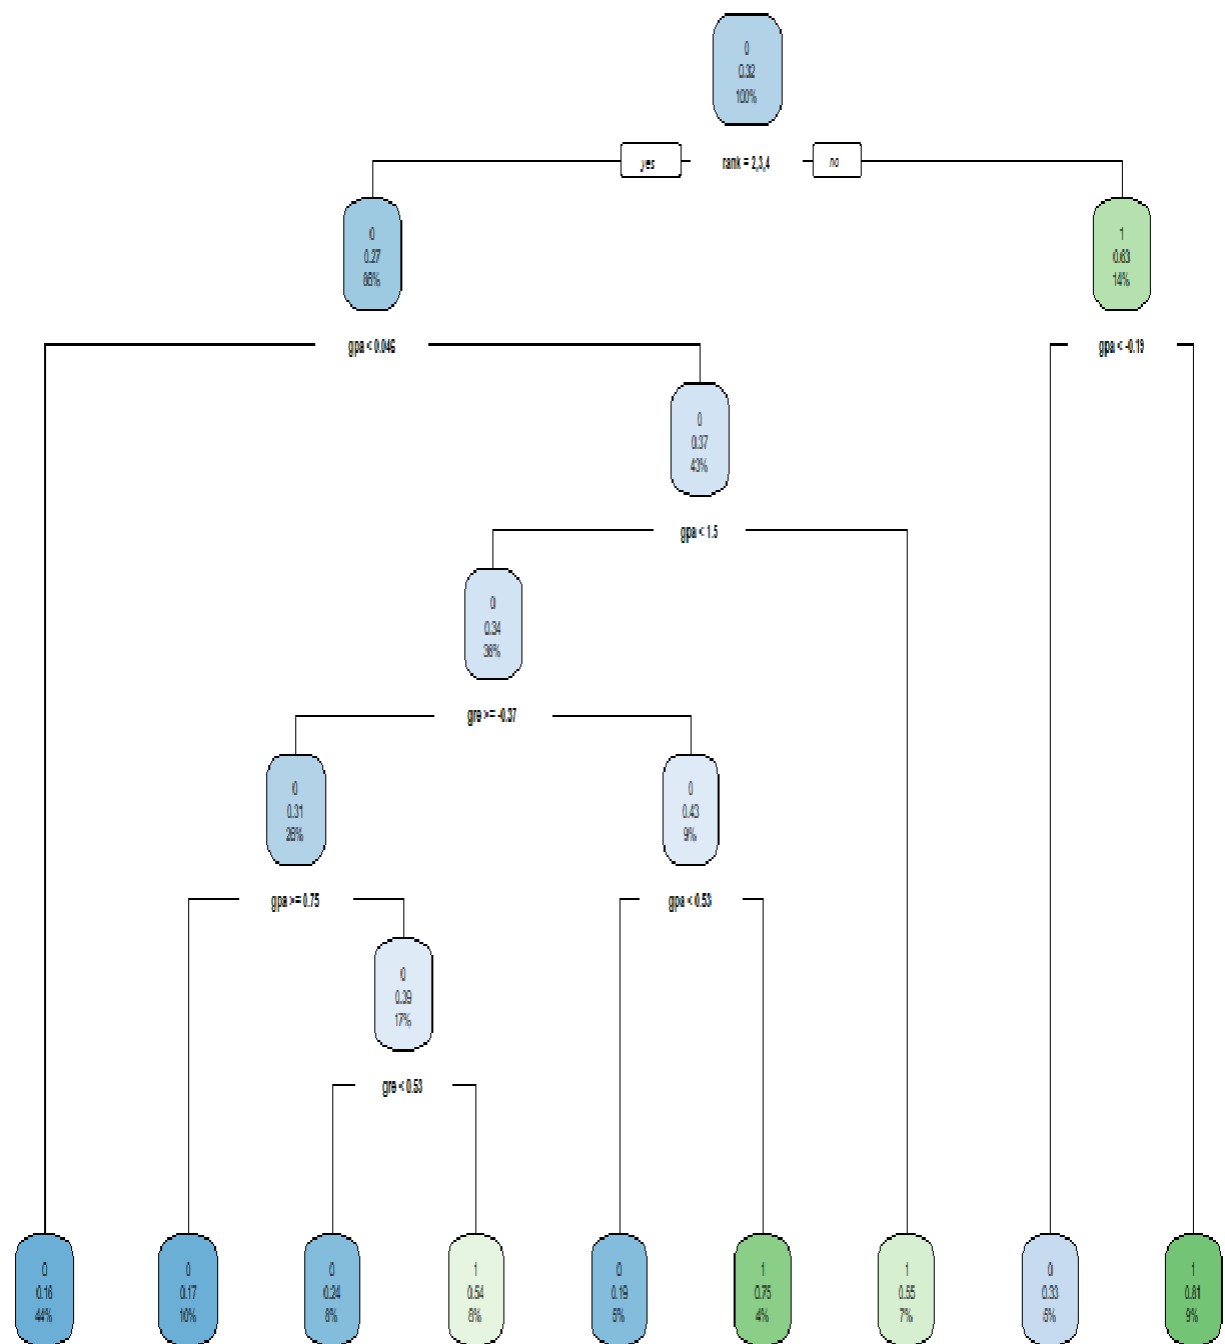
Decision tree:

```
> #Decision tree
> library('rpart')
> library('rpart.plot')
> nrow(train)
[1] 296
> nrow(test)
[1] 99
> 0.03*nrow(train)
[1] 8.88
> 0.03*nrow(test)*3
[1] 8.91
> r.cnt1=rpart.control(minsplit = 26,minbucket = 9,xval=5)
> dec_clf=rpart(admit~.,control=r.cnt1,data=train)
> summary(dec_clf)
Call:
rpart(formula = admit ~ ., data = train, control = r.cnt1)
n= 296
```

	CP	nsplit	rel error	xerror	xstd
1	0.11702128	0	1.0000000	1.0000000	0.08520516
2	0.05319149	1	0.8829787	0.9787234	0.08471012
3	0.02127660	2	0.8297872	0.9148936	0.08309697
4	0.01063830	6	0.7446809	0.9042553	0.08280889
5	0.01000000	8	0.7234043	0.9148936	0.08309697

Variable importance				
gpa	rank	gre	Race	Gender_Male
46	28	17	4	3
ses				
2				

```
Node number 1: 296 observations, complexity param=0.1170213
predicted class=0 expected loss=0.3175676 P(node) =1
class counts: 202 94
probabilities: 0.682 0.318
left son=2 (255 obs) right son=3 (41 obs)
Primary splits:
```



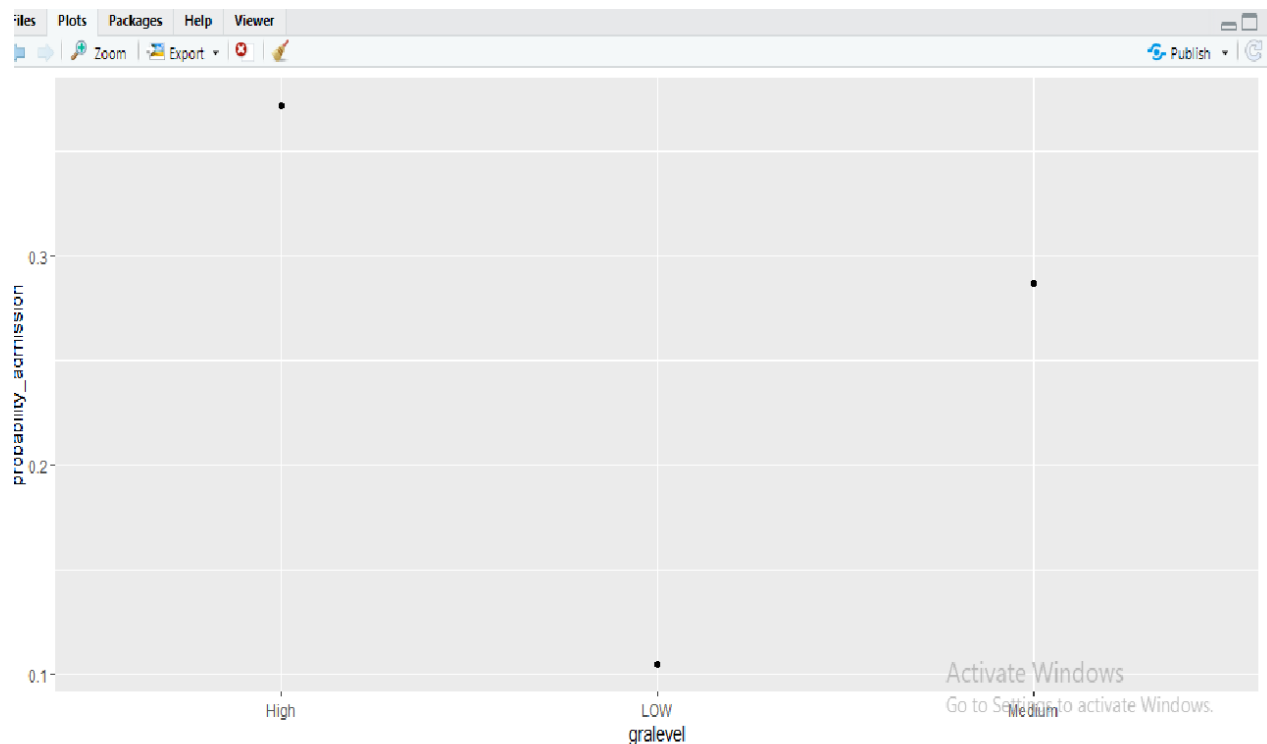
Accuracy of Decision tree:

```
> #Accuracy of Decision tree
>
> predicted_val3=predict(dec_clf,test[-1],type='class')
> predicted_val3
 1  11  14  16  26  29  31  35  37  38  52  60  61  63  68  70  74  82  94  95 102 104 109 113 114 116 118
1  1  0  0  1  0  1  0  0  0  0  0  0  0  0  0  1  1  0  0  1  0  1  0  0  0  1  0
121 126 127 137 138 139 144 150 151 159 161 172 176 179 192 197 198 202 203 205 206 212 214 215 216 223 233
1  0  1  0  1  0  0  0  1  1  0  0  0  0  1  0  0  0  1  1  0  0  0  0  0  0  0  0
236 238 243 247 248 251 253 256 260 266 269 270 274 276 277 279 285 286 288 294 296 304 312 315 317 320 321
0  1  0  0  0  0  1  0  0  0  1  0  0  0  1  0  0  0  0  0  1  0  1  1  0  0  0  0
324 325 332 339 342 358 359 369 370 373 375 382 385 386 391 394 396 400
0  0  1  0  0  0  0  1  0  0  0  0  0  0  0  0  0  1  0
Levels: 0 1
>
> # Confusion matrix
>
> conf_mat3=table(predicted=predicted_val3,actual=test$admit)
> conf_mat3
      actual
predicted 0  1
         0 50 22
         1 17 10
>
> # Accuracy
>
> accuracy3=sum(diag(conf_mat3))/sum(conf_mat3)
> accuracy3
[1] 0.6060606
>
> # Accuracy is 0.6060
>
> ## It is quite evident that Regression model and SVM are the best model with
> # accuracy of 61.61%
> |
```

→ From the above 3 models it is quite evident that Regression and SVM are the best two models with accuracy of 61%.

Categorize the grade points average into High, Medium, Low(with admissionprobability percentage) and plot it into a point chart.

```
> #Categorize the grade point average into High, Medium and Low
>
> categorize=transform(col_adm,gralevel=ifelse(col_adm$gre<440,'LOW',ifelse(col_adm$gre<580,'Medium','High')))
> View(categorize)
>
> sum_descp=aggregate(admit~gralevel,categorize,FUN=sum)
> length_descp=aggregate(admit~gralevel,categorize,FUN=length)
> length_descp
  gralevel  admit
1      High    226
2      LOW     38
3    Medium   136
> probability_table=cbind(sum_descp,Recs=length_descp[,2])
> probability_table
  gralevel  admit  Recs
1      High    84    226
2      LOW     4     38
3    Medium    39   136
> probability_table_final=transform(probability_table,probability_admission=admit/Recs)
> probability_table_final
  gralevel  admit  Recs  probability_admission
1      High    84    226             0.3716814
2      LOW     4     38             0.1052632
3    Medium    39   136             0.2867647
>
> #Plotting the grade point in chart
>
> library(ggplot2)
> ggplot(probability_table_final,aes(x=gralevel,y=probability_admission))+geom_point()
> |
```



	admit	gre	gpa	ses	Gender_Male	Race	rank	gralevel
1	0	380	3.61	1	0	3	3	LOW
2	1	660	3.67	2	0	2	3	High
3	1	800	4.00	2	0	2	1	High
4	1	640	3.19	1	1	2	4	High
5	0	520	2.93	3	1	2	4	Medium
6	1	760	3.00	2	1	1	2	High
7	1	560	2.98	2	1	2	1	Medium
8	0	400	3.08	2	0	2	2	LOW
9	1	540	3.39	1	1	1	3	Medium
10	0	700	3.92	1	0	2	2	High
11	0	800	4.00	1	1	1	4	High
12	0	440	3.22	3	0	2	1	Medium
13	1	760	4.00	3	1	2	1	High
14	0	700	3.08	2	0	2	2	High
15	1	700	4.00	2	1	1	1	High
16	0	480	3.44	3	0	1	3	Medium
17	0	780	3.87	2	0	3	4	High
18	0	360	2.56	3	1	3	3	LOW
19	0	800	3.75	1	1	3	2	High
20	1	540	3.81	1	0	3	1	Medium
21	0	500	3.17	3	0	2	3	Medium
22	1	660	3.67	2	0	2	3	High